

A few stickleback suffice to transport adaptive alleles to new lakes

Jared Galloway, William A. Cresko, and Peter Ralph

December 10, 2018

Abstract

Threespine stickleback fish provide a striking example of local adaptation despite recurrent gene flow. The species is distributed around the Northern hemisphere in both marine and freshwater habitats. It is thought that these numerous, smaller freshwater populations have been established "de novo" from marine fish, and that a shared freshwater phenotype is often established using standing genetic variation. Here we use genealogical simulations to determine the levels of gene flow that best matches observed patterns of allele sharing among habitats in stickleback, and more generally to better understand how gene flow and local adaptation in large metapopulations determine speed of adaptation and reuse of standing genetic variation. (1) We find that rapid, repeated adaptation using a shared set of alleles maintained at low frequency by migration-selection balance occurs over a realistic range of intermediate rates of gene flow. (2) Low gene flow leads to slow, independent adaptation of distinct habitats, whereas high gene flow leads to large migration load. (3) We quantify how F_{ST} scans for adaptive alleles are more likely to succeed with higher rates of gene flow. (4) In addition, we find the origin of many freshwater adapted alleles in the introduced lakes to be propagated from the original generation of marine individuals that had inhabited the lake. The results support existing theory of local adaptation, and provide a more concrete look at a particular, empirically motivated example.

Introduction

The canonical model for the genetics of adaptation, first formulated by Fisher in the early part of the 20th century, involves the sequential fixation of new mutations. While it has proven valid in numerous studies in the field and the lab, this model is now rightfully understood as incomplete for many species in nature that have more complicated population structures. A growing number of studies have identified instances of convergent or parallel adaptive evolution that utilize standing genetic variation. What evolutionary processes promote the maintenance of re-use of such genetic polymorphisms? How do these processes interact with diverse geographic contexts, as well as variation in genetic processes such as recombination, to change the likelihood of so-called standing genetic variation (SGV).

Clues come from population genomic studies in the wild that have recently been powered by advances in sequencing technologies. One such organism is the threespine stickleback. This small fish has long been used for studies of behavior, ecology and evolution in the wild, and has more recently become a model for

understanding the genetic basis of adaptive evolution. The ancestral marine form has given rise to millions of independently derived populations in recently de-glaciated regions of the Northern Hemisphere, which have each evolved a similar suite - or 'freshwater syndrome' - of phenotypes despite local differences. This adaptive phenotypic evolution to freshwater can occur in mere millenia, in many cases since the end of the last ice age about 13,000 years ago. Studies of the fossil record also clearly show that invasion of oceanic stickleback and local adaptation in freshwater has occurred repeatedly over millions of years, often involving very similar phenotypic shifts as seen in contemporary populations.

A finding from recent stickleback population genomic studies is that the same genomic regions, and in some cases the same haplotypes, are implicated in patterns of repeated evolution. These studies show that divergent stickleback populations are clearly still connected via bi-directional gene flow. This reuse of standing genetic variation motivated empirical studies to ask whether this phenotypic and genomic evolution could occur in decades. The answer is a clear yes. For example, in 1964 the Great Alaskan Earthquake caused an uplift of Middleton island and in turn, introduced a group of freshwater ponds around the perimeter of the island. Quickly inhabited by the surrounding marine population of stickleback, Lescak et al. [2015] observed significant phenotypic changes in less than 50 years that appear to be parallel to freshwater stickleback that have been separated for over thousands of years. In freshwater stickleback, the number of lateral plates are reduced and the opercle shapes shows the same expansion of the dorsal region and reduction of the ventral region.

But how can evolution occur so quickly? The leading hypothesis for stickleback is that rather than acting on new mutations, adaptation to freshwater environments is accelerated by selection on SGV found in marine populations. The first clear example of the global reuse of SGV was the gene *eda* which has been shown to be an important regulator for the number of lateral plates. While the low lateral plate version of this gene arose millions of years ago, it is found in freshwater ponds which have formed much more recently. More recently, population genomic studies employing genome-wide haplotype analyses has provided evidence that most regions of the genome that distinguish marine-freshwater genetic differences share this pattern [Nelson and Cresko, 2017]. Genomic regions that exhibit local adaptation, as evidenced by increased relative population genomic divergence, also show much older haplotypes that are enriched in the alternative habitats.

These empirical data generally support the "transporter"-hypothesis proposed by Conte and Schluter in 2009 Schluter and Conte [2009], which is a conceptual model for the flow of freshwater alleles from multiple smaller freshwater populations into much larger and less structured marine populations through hybridization events. Alleles in this asymmetrically structured metapopulation can then be recycled for subsequent freshwater adaptation. Several questions remain, however, about the manner and degree to which these alleles are scattered among the marine fish and must be re-assembled. Is it more akin to the atomization and rebuilding that occurred in the Star Trek series that motivated this hypothesis? Alternatively does it occur in more of a patchwork by individuals, or their early generation hybrids, in geographically adjacent freshwater habitats moving quickly through the marine environment to seed new freshwater populations. More generally, how do different levels of gene flow and local adaptation interact with the asymmetric nature of stickleback population structure affect the dynamics of the origin and rate of adaptive alleles, as well as their eventual organization into marine and freshwater typical genomes?

Here we develop a forward simulation approach in SLiM to model the stickleback evolutionary history in marine and freshwater habitats. We then record the effect of variation in gene flow on the genetic and genomic

architecture of local adaptation. We ask, how rapidly can selection act on standing genetic variants at a given value of migration, M ? Furthermore, how might the inference of causal loci through the use of genome wide association studies (GWAS) and F_{st} scans across the genome be skewed by (lack of) introgression between sub-populations. Finally, we trace back through the genealogical history of adaptive variants in new lakes to find the precise origin of those variants. **Jared - My attempt at concisely stating are questions**

We find that only a few stickleback per generation, per lake, suffice to maintain freshwater alleles at an acceptable frequency in marine environments for efficient “transportation” of the freshwater haplotype. The selection on SGV at this level results in rapid local adaptation in ≈ 60 generations, this follows closely to what we observe in nature (such as middleton island, cite this). With lower $M < 1$ individuals per lake per generation we find either the lakes adapt independently by “*de novo*” mutation, or the selection on SGV is slow. (greater than 2000 gens.) This low level of introgression also results in false positive F_{st} peaks when scanning across the genome suggesting introgression plays a key role in inferring causal loci. With higher $M > 10$ individuals per lake per generation we see the impact of migration load, preventing the demes from fully adapting to their selective pressure. When tracing back to the origin of adaptive alleles being *used* in new lakes that had adapted, we find the large majority of the alleles had been “captured” meaning they were present in the individuals that initially colonized the new lake, common in the original lakes, and uncommon in the ocean. This finding is in opposition to the claim that freshwater alleles are transported through subsequent migration, after the lakes have been colonized. **Jared - My attempt at concisely stating the results**

Methods

We explored these questions using forwards-time simulations with explicit genomic representation of a quantitative trait in SLiM [Haller and Messer, 2017, ?]. The details of the model were motivated by current understanding of threespine stickleback history and demography, but remain simplistic in some aspects due to computational constraints. Possibly the most important caveat is that simulated population sizes are much smaller than the census size of the threespine stickleback population (see the Discussion for more on this).

Habitat and geography Each habitat type – marine and freshwater – have a fixed total of 5,000 diploid individuals each. The arrangement of these habitats, depicted in Figure 1, roughly models a set of freshwater habitats along a stretch of coastline. The marine habitat is a continuous, one-dimensional range of 25 units of length, while the freshwater habitat is divided into ten subpopulations (which we call “lakes”), each connected to the marine habitat at regularly spaced intervals (positions $i - 1/2$ for $1 \leq i \leq 25$).

Divergent selection is mediated by a single quantitative trait with different optima in marine and freshwater habitats. This situation roughly models the cumulative effect of the various phenotypes such as armor morphology, body size, craniofacial variation and opercle shape on which divergent selection is thought to act in the two environments. Concretely, the optimal trait values in the marine and freshwater habitats are +10 and -10 respectively, and fitness of a fish with trait value x in a habitat with optimal value x_{opt} is




Figure 1: **Diagram of simulated populations:** a single, continuous, one-dimensional marine habitat (blue) is coupled to randomly mating “lakes” at discrete locations with dark green arrows representing migration patterns. After an initial period of 100K generations with 25 lakes, an additional 25 lakes are added (at the same set of locations) to simulate the appearance of newly accessible freshwater habitats. The marine habitat, and each set of 25 lakes, each contain 5,000 individuals at all times.

determined by a Gaussian kernel with standard deviation 15, i.e.,

$$f(x; x_{\text{opt}}) = \exp \left\{ \frac{1}{2} \left(\frac{x - x_{\text{opt}}}{15} \right)^2 \right\}.$$

Note that only the scale of trait values relative to mutation effect size is relevant. We chose the difference between optima and strength of stabilizing selection in each habitat so that (a) around 10 (diploid, homozygous) mutations were sufficient to move from one optimum to the other, and (b) well-adapted fish from one habitat would have low, but nonzero, fitness in the other habitat.

Genetic architecture of the trait Each individual carries two copies of a linear chromosome of 10^8 loci. Mutations that can affect the trait under selection can occur at rate 10^{-7} per locus per generation in ten “effect regions” of 10^5 loci each, spread evenly along the chromosome. Each mutation in these regions is either additive, completely recessive, or completely dominant (with equal probability), and with an effect sizes chosen randomly from an Exponential distribution with mean $1/2$, either positive or negative with equal probability. Individual trait values are determined additively from the diploid genotypes. Concretely, an individual that is heterozygous and homozygous for mutations at sets of loci H and D respectively has trait value $x = \sum_{i \in H} h_i s_i + \sum_{j \in D} s_j$, where h_i and s_i are the dominance coefficient and the effect size of the mutation at locus i . Subsequent mutations at the same locus replace the previous allele.

Population dynamics We use SLiM to simulate a Wright–Fisher population with nonoverlapping generations and a fixed population size of 5000 diploid individuals in each habitat. Each generation, the two parents of each new offspring are chosen proportional to their fitness (all individuals are hermaphroditic), and the contributing genomes are produced by Poisson recombination with an average of one crossover per chromosome per generation (10^{-7} per locus per generation). Since the total population across *all* lakes is regulated, to keep population sizes roughly constant within each lake, before selection we divide fitness values of each freshwater individual by the mean fitness in their lake, so that the mean fitnesses of all lakes are equal.

Dispersal occurs both locally along the coastline in the marine habitat and between the marine habitat and the lakes, and can be thought of as occurring at the juvenile stage. There is no dispersal directly between lakes. The lake–ocean migration rate is denoted m , which we refer to as the rate of *gene flow* between habitats. Each new individual in the marine habitat has freshwater parents with probability m ; to obtain the pair, a first parent is chosen proportional to fitness, and a mate is chosen from the same lake as the first, also proportional to fitness. The resulting offspring is given a spatial location in the marine habitat at the location of the parent’s lake. Parents for a new marine individual who is not a migrant are chosen similarly (with probability $1 - m$): first, a single parent is chosen proportionally to fitness in the marine habitat, and then a mate is chosen, also proportionally to fitness but re-weighted by a Gaussian function of the distance separating the two, with standard deviation $1/2$. Concretely, if the first parent is marine individual i , then marine individual j is chosen as the mate with probability proportional to $f(x_j) \exp(-2d_{ij}^2)$, where d_{ij} is the distance between the two locations. Finally, each new marine offspring is given a position displaced from the first parent’s position by a random Gaussian distance with mean 0 and standard deviation 0.02, and reflected to stay within the population. New offspring in the freshwater habitat are chosen in the same way, except the probability that the parents are marine individuals is m ; any new freshwater offspring produced by marine individuals are assigned to the lake nearest to the position of the first marine parent.

New freshwater populations To study how newly appearing freshwater habitats adapt, we introduce a new set of 25 lakes midway through the simulation. The initial set of individuals in these new lakes have parents chosen in the same way from the marine habitat as ocean-to-lake migrants, and act like an independent copy of the original set of lakes – in particular, the two sets of lakes each have 5,000 individuals at all times. (Since this doubles the number of lake-to-marine immigrants, after this happens the probability that a new marine individual has freshwater parents is $2m$ instead of m .)

We quantify genetic differentiation between the habitats with per-locus F_{ST} : if p_f and p_m are the frequencies of a given mutant allele in the freshwater and marine habitats, respectively, and $\bar{p} = (p_f + p_m)/2$, then $F_{ST} = 1 - (p_f(1 - p_f) + p_m(1 - p_m))/(2\bar{p}(1 - \bar{p}))$.

Recording genealogical history We used SLiM’s ability to record *tree sequences* [?] to record the genealogical history of all individuals alive at the time of introduction of new lakes, the time of adaptation, and the end of the simulation. This allowed us to directly query the true origins of adaptive alleles, as well as allowing much larger simulations by avoiding the computationally expensive task of simulating neutral mutations (which were added after the fact at a rate of 10^{-7} per locus per generation, as described in ?).

The output tree sequence from each simulation allows us to, at each locus along the genome, construct the genealogical tree relating all extant chromosomes as well as all chromosomes present at time of introduction of the new lakes. Using these trees, we classified each freshwater allele found in the new lakes at time of adaptation into four categories:

1. *De novo alleles*: deriving from a new mutation that occurred in a new lake.
2. *Migrant alleles*: not present in the individuals that initially colonized the new lake, but brought in by migration at a later time.
3. *Captured alleles*: present in the individuals that initially colonized the new lake, and both common (above 50%) in the original lakes, and uncommon (below 50%) in the ocean.
4. *Marine alleles*: present in the individuals that initially colonized the new lake, and not a captured allele.

I may be conf

Roughly speaking, these four categories describe the contribution of (1) new mutation, (2) post-colonization migration, (3) standing variation at migration–selection balance, and (4) standing variation at mutation–selection balance. We could partition migrant alleles into finer categories (e.g., we do not distinguish whether they are common in the original lakes or not), but this category turns out to be small.

We were also able to use the tree sequence to get information about the individual “effect regions” in all initial genomes of the new lakes at time of introduction. From this we determined the ability for the new lakes to select upon the standing variation in the marine. Given the probability that effect region fixes at a given trait value, $2 * x/45$, where x is the effect size of a haplotype, we calculated the expected total effect sizes of the haplotypes that fix. Concretely, it is the sum across the 10 effect regions of

$$\sum_{i=1}^N \prod_{j < i} (1 - p_j) p_i x_i$$

where N is the number of genomes per lake, p_i is the probability of fixation of the i^{th} haplotype, x_i is the effect size of the i^{th} haplotype, and the haplotypes are sorted in decreasing order by x_i . These numbers were computed assuming additivity of mutations, meaning the numbers produced are a slight overestimate.

Descriptive statistics To assess whether new lakes adapt using existing genetic diversity, we define a *freshwater allele* to be an effect mutation that has frequency higher than 0.5 in at least one of the original lakes, while remaining lower than 0.5 in the marine. This categorization is made for each generation using the allele frequencies from that generation, and so changes with time. Alleles common in the newly introduced lakes do not count if they are not also common in the original lakes. They are defined this way because the transportation hypothesis does not specify where or when an advantageous mutation arises, but simply suggests that any sufficiently common freshwater adapted allele could participate in adaptation in new habitat [Schluter and Conte, 2009].

Time to adaptation of the introduced population, denoted T_{adapt} , is defined to be the generation at which the difference between the average trait value in the original and the introduced freshwater populations is less than 0.5.

Population genetic analyses? Instead of the more general 'Descriptive statistics'?

We describe overall genetic differentiation between the habitats using F_{ST} , calculated on a per-locus basis. Concretely, if p_f and p_m are the frequencies of a given mutant allele in the freshwater and marine habitats, respectively, and $\bar{p} = (p_f + p_m)/2$, then we compute F_{ST} for that mutation as $1 - p_f p_m / (\bar{p}(1 - \bar{p}))$.

Results

We varied the ocean–lake migration rate, m , across separate simulations from 5×10^{-5} to 5×10^{-1} . Since each lake contains 200 individuals, this corresponds to between 0.01 and 100 migrants per lake per generation. Many aspects of adaptation changed substantially across this range, including the speed of adaptation, degree of sharing of adaptive alleles between lakes, and the population genetic signals left behind. At very low rates of gene flow, each new lake’s population adapted *de novo*, which took a very long time. At very high rates of gene flow, local adaptation was almost impossible. Between these two extremes, genetic variation that allowed adaptation to freshwater habitats could move relatively easily between lakes – despite being deleterious in the intervening marine habitat – which allowed populations arriving in new lakes to adapt quickly, reassembling collections of alleles responsible for freshwater adaptation in other lakes.

Local Adaptation: differentiation with gene flow

Local adaptation occurred at most migration rates, as shown in Figure 3. At the highest migration rate (at which half of each population was composed of migrants), populations had only slightly different average trait values. At lower migration rates (10 migrants per generation and below), populations adapted to local conditions: as shown in Figures 2, ?? and ??, freshwater and marine populations diverged until the trait means were close to the optimal values in each habitat. The establishment of new alleles in the lakes is visible in Figure 2 as jumps in the mean trait value, which move the trait by an amount of order 1 every few hundred generations. Trait variation within each population was small compared to the difference between populations. Across all parameter values, differences at around 16 commonly polymorphic sites (eight that shift the trait in each direction) were responsible for most of the adaptive differences between freshwater and marine habitats.

As expected, increasing migration rate decreased differentiation between habitats. As seen in Figure 4, F_{ST} between marine and freshwater habitats at neutral sites steadily declines as migration increases. However, local adaptation was still able to occur despite overall homogenization at up to 10 migrants per lake per generation.

Speed of adaptation Adaptation occurred much more quickly at higher migration rates, both in the old and new sets of lakes. We measured this “time to adaptation” as the number of generations until average trait values in old and new lakes were within 0.5 of each other, shown in Figure 5 for different rates of gene flow. Adaptation of new lakes took over 18,000 generations at the lowest rate, while at one migrant per lake per generation new lakes managed to adapt in only 32 generations.




Figure 2: Mean individual trait values in the marine habitat (blue line), the original lakes (yellow lines; average in orange), and the new lakes (light green lines; average in dark green), across the course of two simulations, with migration rates of **(top)** $m = 5 \times 10^{-5}$ and **(bottom)** $m = 5 \times 10^{-4}$ per generation per individual (i.e., 0.01 and 0.1 migrants per lake per generation, respectively). Optimal trait values in the two habitats are at ± 10 . Analogous plots for other migration rates are shown in Figures ?? and ??.




Figure 3: Distribution of mean trait values across generations of the simulation, for different migration rates. The dashed pink and purple lines at ± 10 give the optimum phenotypes in the marine and freshwater environments, respectively.




Figure 4: Average F_{ST} in windows of 500Mb between: **(top)** marine habitat and old lakes; **(middle)** marine habitat and new lakes; and **(bottom)** old and new lakes. Each plot shows F_{ST} values for a separate simulation, with columns corresponding to increasing gene flow from left to right. Ten vertical dotted pink lines in each subplot show regions which have the potential to affect phenotype.




Figure 5: Time to adaptation as a function of migration rate. The time to adaptation is measured as the number of generations until the introduced population's mean phenotype comes within 0.5 of the original lakes average phenotype. Each point represents a single simulation run, and the yellow dashed line is the average of all points at each respective parameter value. (Adaptation did not occur at the highest rate of gene flow.) **where are the multiple points?, They do not exist, we only ran the one set. They each take ≈ 36 hours to run; but I can run them again if you want.**

Sharing of freshwater adapted alleles

At low migration rates, the *initial* period of adaptation takes roughly ten times longer for lakes as it does for the ocean. This difference occurs because at low migration rates, adaptation occurs independently in each lake, and the marine habitat has ten times as many individuals, and therefore ten times the influx of new alleles than any one lake. At higher migration rates, greater mixing allows lakes to share alleles instead of developing their own genetic basis of adaptation. As a first indication of this, Figure 4 shows that F_{ST} between the “original” and “introduced” sets of lakes at effect mutations decreased with migration rate.

To investigate in more depth how locally adaptive alleles found in the original lakes are shared between lakes, and how they spread to the new lakes, we counted “pre-existing freshwater adapted alleles”. These we define for a particular generation to be any mutation with an effect on the trait whose frequency is above 50% in at least one original lake and below 50% in the marine habitat. Figure 6A shows the distribution of the number of these alleles, across generations. At the lowest migration rate (one migrant per 100 generations per lake), each lake has a private set of about 10 mutations nearly fixed in that lake but not elsewhere: new lakes acquire new adaptive alleles and so have none of these. At one migrant per 10 generations per lake, the original set of lakes each developed private alleles, but these are present at low frequency in the ocean, and so new lakes adapt mostly using a subset of this large repertoire of standing variation. As migration rate increases past this, the number of pre-existing freshwater adapted alleles declines because alleles move between populations by migration before they can appear by new mutation, while the frequency of these alleles in the ocean stays relatively constant.

Figure 6B shows the distribution of the mean percentage of currently-defined freshwater adapted alleles that each genome in each of the populations carries, averaged across time and individuals. If all individuals across lakes carried the same set of alleles determining their trait value, this would be 100%. At the lowest migration rates, each genome in the original lakes have almost exactly $1/25^{th}$ of the total number of pre-existing freshwater adapted alleles – this is because each of the 25 lakes has adapted with a unique set of alleles. Since these are *pre-existing* alleles, the value is zero for introduced lakes. Figure 6A shows us that at 0.1 migrants per lake per generation and above, the average individual across the new lakes has nearly the same amount of pre-existing freshwater adapted alleles as individuals across the old lakes. As expected, the genetic basis of the freshwater phenotype seems to simplify as migration increases – higher rates of migration allow adaptive alleles of higher effect to travel more efficiently through the population, even though they are deleterious in the ocean.

The numbers of Figure 6 suggest that the dramatic increase in speed of local adaptation we observed above occurs because higher gene flow between populations allows sharing of freshwater alleles between populations. We confirmed this by using recorded tree sequences to identify the origin of each trait-affecting allele common in the new lakes at time of adaptation, as defined in the Methods. Figure 7 shows that at the lowest rate of gene flow, the majority of adaptive alleles are derived from *de novo* mutation. As gene flow increases, a larger fraction of adaptive alleles derive from pre-existing variation in the marine population at the time of introduction. In other words, greater mixing at higher migration rates allows lakes to share alleles instead of developing their own genetic basis of adaptation.

At first, increased migration allows sharing of adaptive alleles between lakes, but at the highest migration rate, the constant influx of alleles between the habitats creates substantial migration load. The rate at which migration load becomes substantial, 10 migrants per lake per generation, only replaces 5% of each population

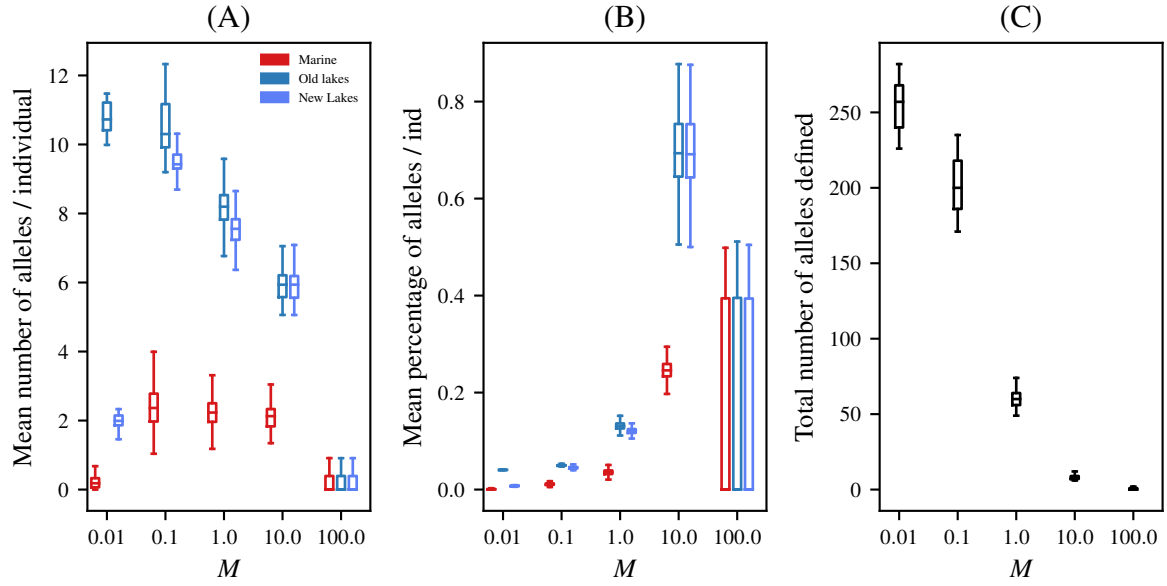


Figure 6: Amount of standing freshwater variation by habitat, across migration rates. Each plot counts “pre-existing freshwater adapted alleles”, that are common in the original lakes but rare in the ocean (see text for definition). **(A)** Mean number of these alleles per individual. **(B)** Mean percentage of these alleles per individual. **(C)** Total number of these alleles (so, $B = A/C$). The number of alleles meeting these conditions changes over the course of the simulation, and each plot shows distributions of these values across generations. The horizontal axis shows M , the mean number of migrants per lake per generation.




Figure 7: **(Origin of adaptive alleles:)** Each bar plot shows the origins of all trait-affecting alleles above frequency 50% in at least one new lake, classified as **(red)** new mutations, **(blue)** post-colonization migrants, **(green)** “captured” from pre-existing lakes, or **(orange)** standing marine variation. See Methods for precise definitions of these categories.

each generation with migrants from the other habitat, but this is sufficient to shift the mean trait values to nearly half their optimal values, as seen in Figure 3.

Standing genetic variation Despite a dramatic difference in the amount of allelic sharing between lakes, standing genetic variance (Figure ??) was around 0.05 in the freshwater habitat across all but the highest migration rates. Concurrently, genetic variation in the marine population steadily increased with migration to a similar value. On the one hand, it is not surprising that lakes, as a group, carry more genetic variation than the marine habitat, since population subdivision allows different alleles to become common in each. However, it is striking that at $m = .005$, the marine habitat carries as much genetic variation, despite a lack of any substantial migration load.

Realized genetic architecture

Now we take a closer look at the genomic architecture of local adaptation between the two habitats. Do the alleles underlying trait differences cluster along the genome? Do measures of local differentiation identify the causal loci? Figure 4 shows plots along the genome of per-locus F_{ST} values between habitats. (Note that we are pooling freshwater habitats; a single lake would provide substantially less power.) Higher migration rates showed more distinct F_{ST} peaks over polymorphic loci underlying trait differences between the habitats. “Background” levels of F_{ST} increase as gene flow decreases, swamping out this signal until the regions under selection are indistinguishable. This is likely due to two reasons: first, stronger genetic drift with less migration leads to higher background F_{ST} , and second, greater sharing of adaptive alleles providing a shared signal across populations.

This suggests that genome scans for local adaptation based purely on measures of differentiation will only be successful given enough migration between habitats. To quantify this, Figure ?? shows the power and false positive rates that would be obtained by an F_{ST} cutoff that declared everything above a certain value to be a causal locus.

In regions of the genome underlying individual trait value, we observed Given that migration increases the gene flow between subpopulations, how valid are F_{st} peaks at different M . Knowing exactly which mutations effect phenotype in our simulations, we can look at the statistical power and false positives given F_{st} per SNP across the genome. In Figure ?? , looking at an F_{st} threshold greater than 1, we see the two lowest migration rates 10^{-5} and 10^{-4} having very little statical power. This along with low false positive rate across all F_{st} threshold values is fairly predictable when you consider the high F_{st} values across the genome.

Origin of introduced freshwater adaptive alleles

We have thus far found that the speed of adaptation depends strongly on the degree to which alleles can be shared between populations. However, there remains a mystery. Figure 6A shows that with $m = 0.1$ migrants per lake per generation, there is substantial sharing of alleles between populations, and yet Figure 5 shows that new lakes take around 2,000 generations to adapt. This is surprising because we see a similar quantity of allele sharing at $m = 1.0$, but much more rapid adaptation (≈ 30 generations). Perhaps this

perhaps we want to remove this bit? revisit if it is telling us something important.

need to know where the FAA are in these figures

Need to fix caption on figure ?? before editing this bit.

60 \times discrepancy occurs because new lakes must wait for an additional influx of freshwater alleles through migration beyond what is present in the initial generation? However, Figure 7 shows that at both $m = 0.1$ and $m = 1.0$, the vast majority of alleles underlying the local adaptation derived from individuals in the original generation (“captured” or “pre-existing”).

An alternative explanation is that at lower rates of gene flow, freshwater alleles present as standing variation in the marine habitat are more tightly linked to marine alleles, and so adaptation in the new lakes must wait for recombination. This might be expected since freshwater alleles present in the ocean, to remain present at migration-selection balance at low migration rates, must be “masked” by nearby compensatory alleles. Reversing this “masking” then slows the process of adaptation that uses this genetic variation. Said another way, higher gene flow from freshwater to the ocean maintains relatively intact freshwater haplotypes that can be more easily rebuilt. Recall that trait-affecting mutations only occur in relatively small regions of 10^5 loci in which recombination occurs at rate 10^{-2} . Even if there is a sufficient amount of variation in the initial population of a lake to shift the trait from the marine optimum (+10) to the freshwater optimum (−10), these may be tightly linked to alleles that counteract their effect. For example, suppose there are 10 variants segregating at low frequency with effect size −1 each, but each is paired with a compensatory allele with effect size +1. Each local haplotype is therefore neutral, and so likely to be at higher frequency in the marine population.

Suppose that the alleles present in the initial population of a new lake have effect sizes (x_1, \dots, x_n) , and that the probability of fixation of a single allele of effect size x is $p(x)$. Let X_i be the average per capita effect of the i^{th} allele after fixation, so that $X_i = 0$ if the allele is lost, and $X_i = 2x_i$ if the allele fixes (ignoring dominance). Therefore, the expected contribution to the final trait value of allele i is $\mathbb{E}[X_i] = 2x_i p(x_i)$. The trait value after the initial round of adaptation is $\sum_i X_i$, and so the expected trait value is $2 \sum_i x_i p(x_i)$. Since there are many alleles that contribute more-or-less additively to the trait, the probability of fixation of a particular allele is not easy to find, but we can roughly approximate it as follows: if the current population trait mean z , then an allele that moves the trait a distance x closer to the optimum has a fitness advantage in heterozygotes of $\exp(-\beta((z-x)^2)/\exp(-\beta z^2)) \approx 1 + 2\beta z x$, where in our parameterization, $\beta = 1/450$. Since the probability of fixation of a single allele with relative advantage s is approximately $2s$ [??], then with $z = 10$ we can use $p(x) \approx 4x/45$. Putting these considerations together, a rough estimate for the total shift in trait value resulting from selection on standing variation in the initial population is $R = \sum_i 8x_i/45$, where the sum is over all trait-affecting alleles. This approximation should be good if each trait-affecting allele is rare (present in only a few copies), effects are additive, and most are unlinked.

We computed R to confirm our suspicion that linkage slowed down adaptation by showing that there is sufficient allelic variation in the initial populations to complete adaptation across both migration rates (0.1 and 1 migrants per lake per generations), but not if the haplotypes of each genomic region are assumed to not recombine.

We calculated for each individual in the initial generation of the new lakes the total contribution to trait values of each of the ten genomic regions that containing trait-affecting mutations. The recombination rate within each of these ten regions is 10^{-2} , so these regions act initially more-or-less like single loci. Even though the trait is additive across loci, the presence of sufficiently many distinct alleles to attain, in principle, the optimal freshwater trait value does not ensure adaptation will actually occur with these alleles. On the contrary, many of these will likely be lost to drift.

To get a rough idea of whether the genetic variation present in the initial population is sufficient for the population to adapt, we translated the total effect size x of each region to a probability of fixation (as described below $p(x) \approx 2x/45$), and then calculated $\sum_i x_i p(x_i)$. If a region moves the trait an average of x units towards the freshwater optimum (averaging over dominance relationships), Given the probability that an “effect region” fixes at a given trait value, $2 \times x/45$ where x is the effect size of a haplotype, we were able to determine expected total effect sizes of the haplotypes that fix (detail about computation in methods). Below, we see the mean of these expected total effect sizes across all lakes.

	m=0.01	m=0.1	m=1	m=10
best	2.77	-12.02	-21.02	-7.02
expected	4.22	-4.78	-14.64	-5.47

Table 1: Haplotypic variation present in the new lakes at time of colonization, across rates of gene flow. **“Best”** shows the most negative trait value achievable with intact haplotypes, averaged across populations. **“Expected”** shows the trait value achieved by averaging haplotypic values weighted by fitness (see text for details). This suggests that populations at $m = 1$ have sufficient genetic variation to completely adapt (reach -10) using intact haplotypes, while at $m = 0.1$ variation must be released from haplotypes by recombination.

Between $m = 0.1$ and $m = 1.0$ we see a large difference in the total effect sizes of the haplotypes that are expected to fix. Given these numbers were computed assuming additivity (slight overestimate) and individuals are diploid, at $m = 0.1$ the mean total genotype expected to fix in the populations puts individuals at a trait value of ≈ -8 after selection on alleles haplotypes in the initial population. In turn, the population must wait for novel mutation or subsequent migration to provide the alleles for the remaining 2 units before attaining a mean average trait value at the optimum (-10). It follows in Figure 2B that the majority time to adaptation is spent within 2 units of mean trait value. The initial adaptation on selection of standing variation up to that point happens quite rapidly. This provides evidence that linkage to marine alleles has a major impact in the introduced population’s ability to select and rebuild a pre-existing haplotype given the standing variation at time of introduction.

Theoretical expectations

Here’s some rough calculations to get a sense for what should be going on. Everything is done in more detail in OTHER PLACES WE SHOULD CITE. Fisher, Chevin, etc.

Suppose a new allele enters a lake, either by migration or mutation. If, when it is rare but present in n copies, it has fitness advantage s – i.e., the expected number of copies in the next generation is $(1 + s)n$ – then the probability that it escapes demographic stochasticity to become common in the population is approximately $2s$ [??]. If the current population all differed from the optimum trait by z , and the allele has effect size $-u$ in heterozygotes, then the fitness advantage of the allele would be $s(u) = \exp(-\beta((z - u)^2) / \exp(-\beta z^2)) \approx 2\beta z u$, where in our parameterization, $\beta = 1/450$. This tells us two things: (1) the rate of adaptation decreases as the population approaches the optimum, and (2) larger mutations (in the right direction) are more likely to fix.

does this
go first or
second?

New mutations The total rate of appearance of new mutations per lake is $\mu_L = 0.04$, which are divided evenly in seven categories: neutral, and then additive, dominant, and recessive in either direction. This implies that a new additive or dominant effect mutation appears once every 87.5 generations, on average. Effect sizes are randomly drawn from an Exponential distribution with mean $1/2$, and so the probability that a dominant mutation manages to establish in a population differing from the optimum by z is roughly $\int_0^\infty 4\beta z u \exp(-2u) du = \beta z$, and so the rate of establishment of dominant mutations is $\beta z/87.5$, i.e., about one such mutation every $2461/z$ generations. The distribution of these successfully established mutations has density proportional to $u \exp(-2u)$, i.e., is Gamma with mean 1 and shape parameter 2. Since additive alleles have half the effect in heterozygotes, they have half the probability of establishment. During the initial phase of adaptation, the populations begin at around distance $z = 10$ from the optimum. Combining these facts, we expect adaptive alleles to appear through mutation at first on a time scale of 250 generations, with the time between local fixation of new alleles increasing as adaptation progresses, and each to move the trait by a distance of order 1.

Standing variation An allele that moves the trait z units in the freshwater direction in heterozygotes has fitness roughly $\exp(-\beta z^2) \approx 1 - \beta z^2$ in the marine environment (which is close to optimal). The product of population size and fitness differential in the marine environment for a mutation with $z = 1$ is therefore $2Ns = 8.9$, implying that these alleles are strongly selected against but might occasionally drift to moderate frequency. The average frequency of such an allele in the marine environment at migration-selection equilibrium is equal to the total influx of alleles per generation divided by the selective disadvantage, which if $M = 2000m$ is the number of immigrants per generation, is $2M/\beta z^2$. With $z = 1$, the factor multiplying M is $2/\beta z^2 \approx 1/200$: since lakes have 400 genomes, as long as $M \geq 1$, the chances are good that any particular lake-adapted allele that is present in all pre-existing lakes will appear at least once in the fish that colonize a new lake. However, an allele with $z = 1$ only has probability of around $1/20$ of establishing locally, suggesting that we'd need $M \geq 10$ to ensure enough pre-existing genetic variation that adaptation would happen entirely using the initial set of colonizers. This corresponds to our highest two migration rates, as in e.g., Figure 2B.

Migration If sufficient genetic variation is not present in a new lake initially, it must appear by new mutation or migration. Since a proportion m of each lake is composed of migrants each generation, it takes $1/m$ generations until the genetic variation introduced by migrants equals the amount initially present at colonization. This implies a dichotomy: either (a) adaptation is possible using variants present at colonization or arriving shortly thereafter, or (b) adaptation takes many multiples of $1/m$ generations.

These calculations depend on there being no bottleneck in colonization of the lake; if there is a bottleneck, then an additional factor must be added. At what point do we expect new mutation to be more important than migration for adaptation? By the calculations above, if $M \geq 10$, we expect initial diversity in a lake to be sufficient for adaptation, corresponding to our third-highest migration rate. If this does not happen, then we expect adaptation to take a multiple of $1/m$ generations; with our values, $1/m$ ranges from 20,000 to 20 generations. Above we estimated that adaptive alleles due to new mutation fix locally about every 2,000 generations, suggesting that at our second-lowest migration rate (where $1/m = 2,000$), the two contributions of migration and new mutation are roughly equal. This is in fact what we see: in Figure 6, we see that at

the second migration rate, alleles start to be shared between lakes, while by the third migration rate, they are almost entirely shared.

Discussion

We have shown that historical introgression, at our given parameter sets, is able to reproduce rapid and parallel adaptation similar to what we've seen in real populations such as Middleton island. Selection is able to rebuild the freshwater haplotype from marine populations as a medium between all freshwater populations. Almost all rates of migration were helpful in the efficiency of the population to locally adapt except for the highest at which migration load limited the ability of the populations to reach the local optimum.

We also have also explored the genomic architecture as a consequence of the nature of selection in our scenario. The alleles underlying individual trait value were of large effect and a low number. With a total of 10^5 loci, to be realistic, each loci should represent 1000 *Kb* in real data.

We have also shown introgression is beneficial for inferring causative loci from divergence (F_{st}) along the genome. This is generally because noise of selectively neutral alleles divergence can appear causative when genetic drift causes more differences between populations that have little gene flow between them. It's important to know that in all scenarios, hitchhiking of selectively neutral alleles could also be mistaken for being causative as they often display the same amount of divergence.

thresholds

We have found that too little migration leads to selection upon new mutations in all subpopulations and lakes alike. In contrast, at high migration rates we have seen that migration load limits the ability for species to locally adapt to the selective pressure of their environment. This leads us to consider a window of introgression which allows for the transportation of FAA's without migration load.

connect results back to real data?

The adaptive filter? RAMBLING THOUGHTS HERE Since larger effect alleles are more likely to establish, be it by mutation or migration, repeated colonization of new freshwater habitats will select for larger alleles, be it single alleles or haplotypes bound together by an inversion. However, these are more strongly selected against in the interstitial time. Being recessive would help with this, but would also make it more difficult to establish.

This also brings to light the role of recombination in a system of migration selection balance. On one hand, a lower recombination rate would allow for the freshwater haplotype to remain "intact" within the marine environment, however, without significant introgression the haplotype would quickly get selected against in marine populations. On the other hand a higher recombination rate would, in theory, allow freshwater adapted alleles more longevity as they would hitchhike with marine adapted haplotypes.

Modeling assumptions Our simulations had much smaller population sizes than real populations. How might this affect things?

TODO: effect of smaller than realistic population sizes?

TODO: Talk about the role recombination plays

Below may belong in the discussion / future work

References

- Benjamin C. Haller and Philipp W. Messer. Slim 2: Flexible, interactive forward genetic simulations. *Molecular Biology and Evolution*, 34(1):230–240, 2017. doi: 10.1093/molbev/msw211. URL <http://dx.doi.org/10.1093/molbev/msw211>.
- Emily A. Lescak, Susan L. Bassham, Julian Catchen, Ofer Gelmond, Mary L. Sherbick, Frank A. von Hippel, and William A. Cresko. Evolution of stickleback in 50 years on earthquake-uplifted islands. *Proceedings of the National Academy of Sciences*, 112(52):E7204–E7212, 2015. ISSN 0027-8424. doi: 10.1073/pnas.1512020112. URL <http://www.pnas.org/content/112/52/E7204>.
- Thomas C. Nelson and William A. Cresko. Ancient genomic variation underlies recent and repeated ecological adaptation. *bioRxiv*, 2017. doi: 10.1101/167981. URL <https://www.biorxiv.org/content/early/2017/07/25/167981>.
- Peter L. Ralph and Graham Coop. Convergent evolution during local adaptation to patchy landscapes. *PLOS Genetics*, 11(11):1–31, 11 2015. doi: 10.1371/journal.pgen.1005630. URL <https://doi.org/10.1371/journal.pgen.1005630>.
- Dolph Schluter and Gina L. Conte. Genetics and ecological speciation. *Proceedings of the National Academy of Sciences*, 106(Supplement 1):9955–9962, 2009. ISSN 0027-8424. doi: 10.1073/pnas.0901264106. URL http://www.pnas.org/content/106/Supplement_1/9955.

Supplementary material